



The (Mutable) Config Management Showdown

...TL;DR version



Who am I?




Bob Killen / rkillen@umich.edu / @mrbobbytables

Senior Research Cloud Administrator @ ARC-TS

<http://arc-ts.umich.edu>



Configuration Management Landscape





Paring things down: Mutable vs Immutable Infrastructure

Mutable

- Object is Managed in Place.
- Object State can drift.

Immutable

- Object is built before deployment.
- Mutable systems can be used to build an immutable object.



Who are the (mutable) players?





ANSIBLE

Language:
Python

DSL:
yaml

Architecture:
Agentless

Execution:
Imperative



CHEF™

Language:
Ruby

DSL:
-

Architecture:
Agent/Server

Execution:
Imperative



Language:
Ruby

DSL:
PuppetDSL

Architecture:
Agent/Server

Execution:
Declarative



SALTSTACK

Language:
Python

DSL:
yaml*

Architecture:
Agent/Server*

Execution:
Declarative or
Imperative



ANSIBLE

Ansible: Overview

- Created in 2012 by Michael DeHaan
- Designed to be 'radically simple' (KISS)
- Uses INI based file format for host inventory and group management.
- Executes '**playbooks**' consisting of **tasks**, **roles**, or other **playbooks**.
- Secrets managed through '**ansible-vault**'.
- Central Server/API available through Ansible Tower / AWX
- Large Users:
 - Atlassian
 - Capital One
 - Verizon



ANSIBLE

Ansible: Task Examples

```
- name: Create Mitchell
  user:
    name: mitchell
    uid: 1000
    shell: /bin/bash
    home: /home/mitchell
```

```
- name: Start httpd
  service:
    name: httpd
    state: started
    enabled: true
    when: ansible_distribution == 'Ubuntu'
```




ANSIBLE

Ansible: Anatomy of a Run

1. **Host Inventories**, **Playbooks** and **Roles** are stored on server.
2. A **playbook** is called referencing members of the **host inventory**.
3. Ansible server connects to hosts via ssh and copies the needed python scripts (**modules**) to complete the **playbook**.
4. The **playbook** is executed on the host with the results of each task being returned to the server before the next task is executed.



ANSIBLE

Ansible: Pros & Cons

Pros:

- SSH-Based
- Low Barrier of entry
- Supports a variety of devices
- Easy to extend with Modules
- Good aws support

Cons:

- Master must be able to reach target
- DSL is extremely limiting
- Slow with a large number of hosts
- Very little windows support



Chef: Overview

- Created in 2009 by former AWS engineers
- No DSL, fully Ruby based with an erlang powered API server.
- Full rest/json api
- Compliance automation checks built-in
- Executes '**runlists**' containing **roles**, and **recipes**
- Secrets are managed via encrypted '**data-bags**'
- Large Customers:
 - Facebook
 - Riot Games
 - Microsoft

Chef: Resource Example

```
user 'mitchell' do
  uid '1000'
  gid '1000'
  home '/home/mitchell'
  shell '/bin/bash'
end
```

```
if node[:platform] == 'ubuntu'
  service 'httpd' do
    action [ :enable, :start ]
  end
end
```



Chef: Anatomy of a Run

1. Client rebuilds node information via **ohai**
2. Client contacts server and pull down **node object** containing **run list**.
3. Client pulls down **cookbooks** contained in **run-list**.
4. **Node object** is rebuilt and the run is compiled
5. **Run-list** is executed and node is **converged**.
6. Updated **node object** with results is returned to server.



Chef: Pros & Cons

Pros:

- Mature
- Very easy to learn if coming from a programming background.
- Massive amount of community contributions and tools.
- Tightly integrated with AWS Opswork
- Strong Windows Support

Cons:

- Production level deployment is a challenge if not using managed solution.
- Steep learning curve if not familiar with Ruby.
- Secret Management is problematic without 3rd party tools

A horizontal bar with a teal segment on the left and an orange segment on the right.

Puppet: Overview

- Created in 2005 by Luke Kanies.
- json like DSL describing resources and their dependencies.
- Client/Server communicate over TLS
- Full rest/json API
- Generates a DAG to determine what clients execute
- Large Customers:
 - CERN
 - Walmart
 - Wikimedia Foundation

Puppet: Resource Example

```
user { 'mitchell':  
  ensure => present,  
  uid     => '1000',  
  gid     => '1000',  
  shell   => '/bin/bash',  
  home    => '/home/mitchell'  
}
```

```
if $::osfamily == 'Ubuntu' {  
  service { 'httpd':  
    ensure => running,  
    enable => true  
  }  
}
```


A horizontal bar with a teal segment on the left and an orange segment on the right.

Puppet: Anatomy of a Run

1. Client rebuilds **fact index** by running **facter**
2. Client checks in with server.
3. Desired agent state (**catalog**) is calculated on server and produced a DAG
4. **Catalog** is sent to client and is evaluated against actual state
5. Result is returned to server

A horizontal bar with a teal segment on the left and an orange segment on the right.

Puppet: Pros & Cons

Pros:

- Mature
- Very Simple DSL
- Strong Windows Support
- Very good UI
- Strong Reporting Functionality

Cons:

- Difficult to troubleshoot
- Customizing it requires Ruby knowledge
- Can be complicated to scale
- Can be slow



SALTSTACK

Salt: Overview

- Created in 2011
- 'Event-Driven' orchestration, designed to execute tasks quickly and in parallel
- 'Grain' or fact driven inventory system
- Mainly uses yaml and jinja, but supports many other templating engines and data sources.
- Full json/rest API
- Large User:
 - Cloudflare
 - LinkedIn
 - Lyft



Salt: State Examples

Create Mitchell:

user.present:

- name: mitchell
- uid: 1000
- gid: 1000
- shell: /bin/bash
- home: /home/mitchell

```
{% if grains['os'] == 'Ubuntu' %}
```

Start httpd:

service.running:

- name: httpd
- enabled: true

```
{% endif %}
```

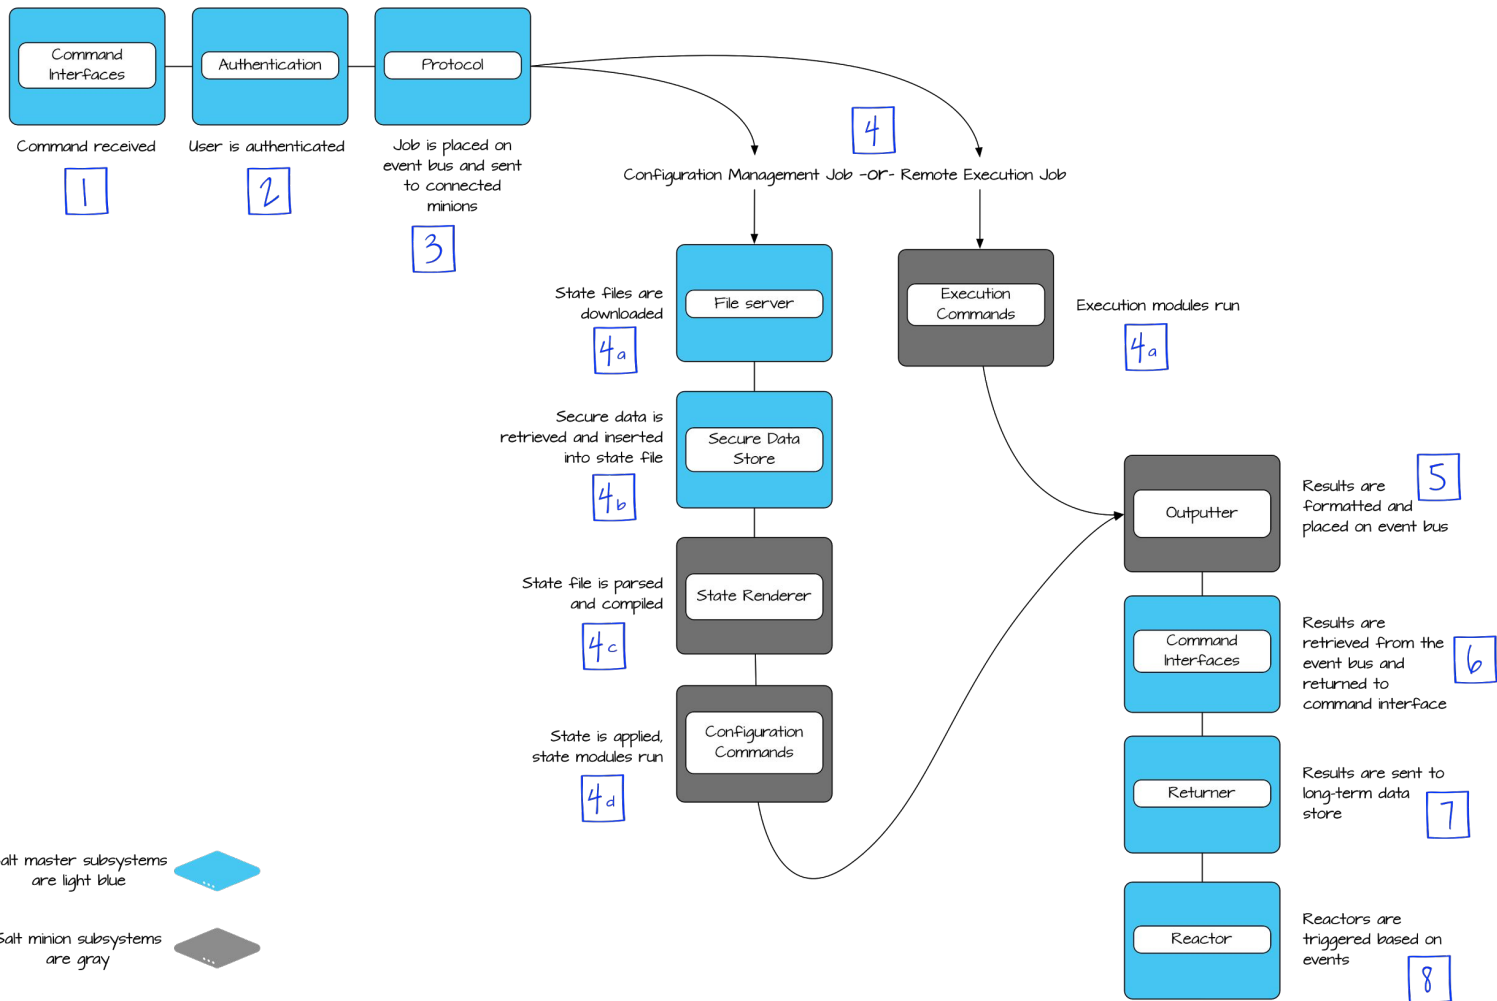


SALTSTACK



Salt: Anatomy of a Run

...it's complicated





Salt: Pros & Cons

Pros:

- Extremely flexible
- Powerful inventory system
- Built in event-engine
- Highly scalable architecture
- Python 3.x support

Cons:

- Exceedingly high learning curve
- Difficult to debug
- Event-bus nature of the system can make it hard to secure in a multi-tenant fashion



Wrapping up

- Use Ansible when...
 - Need a CM system up and going yesterday
 - Your inventory and applications are fairly static and don't require much in the way of conditional cases
- Use Chef if...
 - You're coming from a programming background
 - Tight cloud integration is a boon
- Use Puppet when...
 - Have little to no programming experience
 - UI and reporting are important
- Use Salt if...
 - You have complex orchestration tasks
 - Event driven infrastructure makes sense for your organization